

Reg No : 23MCA1030

Name : Vinayak Kumar Singh

**Exercise 10 : Hierarchical clustering**Collab url : <https://colab.research.google.com/drive/1HK-FQPz1ngNORohkjjWMT07eEUPSaMD1?usp=sharing>

## Import necessary libraries

```
import pandas as pd
import numpy as np
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree
from sklearn.ensemble import RandomForestClassifier
```

## Load the dataset

```
data = pd.read_csv('/content/StudentsPerformance.csv')
data.info();
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   gender                                1000 non-null   object
1   race                                  1000 non-null   object
2   parental_education                    1000 non-null   object
3   lunch                                 1000 non-null   object
4   test_preparation_course               1000 non-null   object
5   math_score                            1000 non-null   int64
6   reading_score                         1000 non-null   int64
7   writing_score                          1000 non-null   int64
dtypes: int64(3), object(5)
memory usage: 62.6+ KB
```

data

	gender	race	parental_education	lunch	test_preparation_course	math_score
0	female	group B	bachelor's degree	standard	none	
1	female	group C	some college	standard	completed	
2	female	group B	master's degree	standard	none	
3	male	group A	associate's degree	free/reduced	none	
4	male	group C	some college	standard	none	
...	...	...	...	...	...	
995	female	group E	master's degree	standard	completed	
996	male	group C	high school	free/reduced	none	

Next steps: [Generate code with data](#)[View recommended plots](#)

## Data preprocessing

## Encode categorical variables

```
label_encoders = {}
```

```

for column in ['gender', 'race', 'parental_education', 'lunch', 'test_preparation_course']:
    label_encoders[column] = LabelEncoder()
    data[column] = label_encoders[column].fit_transform(data[column])

# Feature selection
features = ['gender', 'race', 'parental_education', 'lunch', 'test_preparation_course',
            'math_score', 'reading_score', 'writing_score']
X = data[features]
y = data['math_score'] # Predicting math scores, we can change this to reading_score or writing_score

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Model training
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Model evaluation
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error:", mse)
print("R^2 Score:", r2)

```

```

Mean Squared Error: 1.1386555
R^2 Score: 0.9953206912677982

```

```

# Define grading criteria
def get_grade(score):
    if score >= 90:
        return 'A'
    elif score >= 80:
        return 'B'
    elif score >= 70:
        return 'C'
    elif score >= 60:
        return 'D'
    else:
        return 'F'

# Apply grading criteria to predicted scores
predicted_grades = [get_grade(score) for score in y_pred]

# Get actual grades from the test set
actual_grades = [get_grade(score) for score in y_test]

# Compare predicted grades with actual grades
correct_predictions = sum(1 for pred, actual in zip(predicted_grades, actual_grades) if pred == actual)
total_predictions = len(predicted_grades)

# Calculate accuracy
accuracy = correct_predictions / total_predictions
print("Accuracy:", accuracy)

```

```

Accuracy: 0.99

```

```

# Define a function to calculate the score based on groups
def calculate_group_score(group):
    group_data = data[data['race'] == group]
    score = group_data[['math_score', 'reading_score', 'writing_score']].mean(axis=0)
    return score.mean()

# Calculate scores for each group
group_scores = {}
for group in data['race'].unique():
    group_scores[group] = calculate_group_score(group)

# Print the scores for each group
for group, score in group_scores.items():
    print(f"Group {group}: {score}")

```

```

Group 1: 65.46842105263157
Group 2: 67.13166144200626
Group 0: 62.99250936329588
Group 3: 69.1793893129771
Group 4: 72.75238095238096

```

```
# Calculate total score for each student
data['total_score'] = data['math_score'] + data['reading_score'] + data['writing_score']

# Define grading criteria based on total marks
def get_grade(total_score):
    if total_score >= 240:
        return 'A'
    elif total_score >= 180:
        return 'B'
    elif total_score >= 120:
        return 'C'
    else:
        return 'D'

# Assign grades to students based on total marks
data['grade'] = data['total_score'].apply(get_grade)

# Print the grades for each student
print(data[['gender', 'race', 'parental_education', 'lunch', 'test_preparation_course', 'total_score', 'grade']])
```

	gender	race	parental_education	lunch	test_preparation_course	\
0	0	1	1	1		1
1	0	2	4	1		0
2	0	1	3	1		1
3	1	0	0	0		1
4	1	2	4	1		1
..	...	...	...	...		...
995	0	4	3	1		0
996	1	2	2	0		1
997	0	2	2	0		0
998	0	3	4	1		0
999	0	3	4	0		1

	total_score	grade
0	218	B
1	247	A
2	278	A
3	148	C
4	229	B
..	...	...
995	282	A
996	172	C
997	195	B
998	223	B
999	249	A

[1000 rows x 7 columns]

```
# Define grading criteria based on total marks
def get_grade(total_score):
    if total_score >= 240:
        return 'A'
    elif total_score >= 180:
        return 'B'
    elif total_score >= 120:
        return 'C'
    else:
        return 'D'

# Calculate total score for each student
data['total_score'] = data['math_score'] + data['reading_score'] + data['writing_score']

# Assign grades to students based on total marks
data['grade'] = data['total_score'].apply(get_grade)

# Function to get grade based on index number
def get_grade_by_index(index_no):
    student_row = data.iloc[index_no] # Retrieve the row corresponding to the index number
    return student_row['grade'] # Retrieve the grade from the 'grade' column

# Example usage:
index_no = 10 # Example index number
grade = get_grade_by_index(index_no)
print(f"Grade of student at index {index_no}: {grade}")
```

Grade of student at index 10: C

## ✓ Check for missing values in the dataset

```
missing_values = data.isnull().sum()
print("Missing Values:")
print(missing_values)
```

```

Missing Values:
gender                0
race                  0
parental_education    0
lunch                 0
test_preparation_course 0
math_score            0
reading_score         0
writing_score         0
total_score           0
grade                0
dtype: int64

```

```

# Mapping numeric values to categories
data['gender'] = data['gender'].map({0: 'female', 1: 'male'})

```

```

# Check unique values in the 'gender' column after mapping
unique_genders = data['gender'].unique()
print("Unique Genders:", unique_genders)

```

```

Unique Genders: ['female' 'male']

```

```

# Calculate average scores for male and female students
male_avg_scores = data[data['gender'] == 'male'][['math_score', 'reading_score', 'writing_score']].mean()
female_avg_scores = data[data['gender'] == 'female'][['math_score', 'reading_score', 'writing_score']].mean()

```

```

# Print average scores for male and female students
print("Average Scores for Male Students:")
print(male_avg_scores)
print("\nAverage Scores for Female Students:")
print(female_avg_scores)

```

```

Average Scores for Male Students:
math_score      68.728216
reading_score    65.473029
writing_score    63.311203
dtype: float64

```

```

Average Scores for Female Students:
math_score      63.633205
reading_score    72.608108
writing_score    72.467181
dtype: float64

```

## ✎ Hierarchial Clustering

```

from sklearn.cluster import AgglomerativeClustering
import matplotlib.pyplot as plt
from sklearn.metrics import silhouette_score, davies_bouldin_score

```

```

# Define features (math, reading, writing scores)
features = ['math_score', 'reading_score', 'writing_score']
X = data[features]

```

```

# Choose the number of clusters (k)
k = 3 # we can choose any number of clusters based on our requirements

```

```

# Create AgglomerativeClustering model
hierarchical_clustering = AgglomerativeClustering(n_clusters=k)

```

```

# Fit the model to the data
hierarchical_clustering.fit(X)

```

```

# Get cluster labels
cluster_labels = hierarchical_clustering.labels_

```

```

# Add cluster labels to the dataset
data['cluster'] = cluster_labels

```

```

# Visualize the clusters (for 2D representation)
plt.scatter(X['math_score'], X['reading_score'], c=cluster_labels, cmap='viridis')
plt.xlabel('Math Score')
plt.ylabel('Reading Score')
plt.title('Hierarchical Clustering')

```

```

# Print the number of students in each cluster
print("Number of students in each cluster:")
print(data['cluster'].value_counts())

```

```

# Calculate silhouette score
silhouette = silhouette_score(X, cluster_labels)

```

```

# Calculate Davies-Bouldin index

```

```
davies_bouldin = davies_bouldin_score(X, cluster_labels)
```

```
print("Silhouette Score:", silhouette)  
print("Davies-Bouldin Index:", davies_bouldin)
```

```
plt.show()
```




Number of students in each cluster:

```
cluster  
0    534  
2    314  
1    152
```


Name: count, dtype: int64

Silhouette Score: 0.35242898815984935

Davies-Bouldin Index: 0.8487755692543697



```
import matplotlib.pyplot as plt  
from scipy.cluster.hierarchy import dendrogram, linkage  
  
# Define features (math, reading, writing scores)  
features = ['math_score', 'reading_score', 'writing_score']  
X = data[features]  
  
# Choose the number of clusters (k)  
k = 3 # we can choose any number of clusters based on our requirements  
  
# Create AgglomerativeClustering model  
hierarchical_clustering = AgglomerativeClustering(n_clusters=k)  
  
# Fit the model to the data  
hierarchical_clustering.fit(X)  
  
# Get cluster labels  
cluster_labels = hierarchical_clustering.labels_  
  
# Add cluster labels to the dataset  
data['cluster'] = cluster_labels  
  
# Plot the dendrogram  
plt.figure(figsize=(15, 10))  
plt.title('Hierarchical Clustering - Dendrogram')  
plt.xlabel('Data Points')  
plt.ylabel('Distance')  
dendrogram(linkage(X, method='ward'), orientation='top', distance_sort='descending')  
plt.show()
```



✓ Visualize the clusters (for 2D representation)

```
plt.figure(figsize=(10, 7))
plt.scatter(X['math_score'], X['reading_score'], c=cluster_labels, cmap='viridis')
plt.xlabel('Math Score')
plt.ylabel('Reading Score')
plt.title('Hierarchical Clustering - Scatter Plot')
plt.show()
```

